

# IMDB DATASET

In [5]:

```
# How many rows are there in the IMDB dataset?

import pandas as pd
import numpy as np
df = pd.read_csv("IMDB-Movie-Data.xlsx.csv")
print(df)

# What is the 75th percentile of rating in the IMDB dataset?

df.describe()
```

	ID	Title	Genre	Director	Year	\
0	1	Guardians of the Galaxy	Action	James Gunn	2014	
1	2	Prometheus	Adventure	Ridley Scott	2012	
2	3	Split	Horror	M. Night Shyamalan	2016	
3	4	Sing	Animation	Christophe Lourdelet	2016	
4	5	Suicide Squad	Action	David Ayer	2016	
...	...	...	...	...	...	
995	996	Secret in Their Eyes	Crime	Billy Ray	2015	
996	997	Hostel: Part II	Horror	Eli Roth	2007	
997	998	Step Up 2: The Streets	Drama	Jon M. Chu	2008	
998	999	Search Party	Adventure	Scot Armstrong	2014	
999	1000	Nine Lives	Comedy	Barry Sonnenfeld	2016	

	Runtime_minutes	Rating	Votes	Revenue_millions
0	121	8.1	757074	333.13
1	124	7.0	485820	126.46
2	117	7.3	157606	138.12
3	108	7.2	60545	270.32
4	123	6.2	393727	325.02
...	...	...	...	...
995	111	6.2	27585	NaN
996	94	5.5	73152	17.54
997	98	6.2	70699	58.01
998	93	5.6	4881	NaN
999	87	5.3	12435	19.64

[1000 rows x 9 columns]

Out[5]:

	ID	Year	Runtime_minutes	Rating	Votes	Revenue_millions
count	1000.000000	1000.000000	1000.000000	1000.000000	1.000000e+03	872.000000
mean	500.500000	2012.783000	113.172000	6.723200	1.698083e+05	82.956376
std	288.819436	3.205962	18.810908	0.945429	1.887626e+05	103.253540
min	1.000000	2006.000000	66.000000	1.900000	6.100000e+01	0.000000
25%	250.750000	2010.000000	100.000000	6.200000	3.630900e+04	13.270000
50%	500.500000	2014.000000	111.000000	6.800000	1.107990e+05	47.985000
75%	750.250000	2016.000000	123.000000	7.400000	2.399098e+05	113.715000
max	1000.000000	2016.000000	191.000000	9.000000	1.791916e+06	936.630000

In [6]:

```
# How many NA values are there in the field 'Revenue'?

df.isna().value_counts()
```

Out[6]:

```
ID      Title  Genre  Director  Year  Runtime_minutes  Rating  Votes  Revenue_millions
False  False  False  False      False  False          False  False  False
872
128
dtype: int64
```

In [7]:

```
df = pd.read_csv("IMDB-Movie-Data.xlsx.csv")
df
```

Out[7]:

	ID	Title	Genre	Director	Year	Runtime_minutes	Rating	Votes	Revenue_millions
0	1	Guardians of the Galaxy	Action	James Gunn	2014	121	8.1	757074	333.13
1	2	Prometheus	Adventure	Ridley Scott	2012	124	7.0	485820	126.46
2	3	Split	Horror	M. Night Shyamalan	2016	117	7.3	157606	138.12
3	4	Sing	Animation	Christophe Lourdelet	2016	108	7.2	60545	270.32
4	5	Suicide Squad	Action	David Ayer	2016	123	6.2	393727	325.02
...	...	...	...	...	...	...	...	...	...
995	996	Secret in Their Eyes	Crime	Billy Ray	2015	111	6.2	27585	NaN
996	997	Hostel: Part II	Horror	Eli Roth	2007	94	5.5	73152	17.54
997	998	Step Up 2: The Streets	Drama	Jon M. Chu	2008	98	6.2	70699	58.01
998	999	Search Party	Adventure	Scot Armstrong	2014	93	5.6	4881	NaN
999	1000	Nine Lives	Comedy	Barry Sonnenfeld	2016	87	5.3	12435	19.64

1000 rows x 9 columns

In [8]:

```
# How many movies have revenue higher than 75 million?

revenue = 75.000000
filter1 = df[df["Revenue_millions"] > revenue ]
filter2 = filter1.count()
filter2
```

Out[8]:

```
ID      318
Title    318
Genre    318
Director 318
Year     318
Runtime_minutes 318
Rating    318
Votes     318
Revenue_millions 318
dtype: int64
```

In [8]:

```
# How many movies have revenue greater than 50 million but rating less than 7?

revenue = 50.000000
filter3 = df[(df["Revenue_millions"] > revenue) & (df["Rating"] < 7)]
filter4 = filter3.count()
filter4
```

Out[8]:

```
ID      211
```

Title 211  
Genre 211  
Director 211  
Year 211  
Runtime\_minutes 211  
Rating 211  
Votes 211  
Revenue\_millions 211  
dtype: int64

In [9]:

```
# What is the total revenue generated by movies in the year 2015?  
  
x = df[(df["Year"]==2015)]  
x1 = x.iloc[:,8]  
x2 = x1.sum()  
x2
```

Out[9]:

8854.119999999999

In [10]:

```
# What is the average rating for the genre adventure in the year 2015?  
  
df  
x = df[(df["Year"]==2015) & (df["Genre"]=="Adventure")]  
x2 = x["Rating"].mean()  
x2
```

Out[10]:

6.8

In [11]:

```
# What is the average duration of movies in rows 75 to 150? Please note that the rows in  
python start from 0.  
  
x = df.iloc[75:150]  
x  
avg = x["Runtime_minutes"].mean()  
avg
```

Out[11]:

127.61333333333333

In [12]:

```
# Which year generated the highest revenue?  
  
df.groupby(by="Year").sum().sort_values(by="Revenue_millions",ascending=False)  
  
C:\Users\hp\AppData\Local\Temp\ipykernel_5340\3506442192.py:3: FutureWarning: The default  
value of numeric_only in DataFrameGroupBy.sum is deprecated. In a future version, numeric  
_only will default to False. Either specify numeric_only or select only columns which sho  
uld be valid for the function.  
df.groupby(by="Year").sum().sort_values(by="Revenue_millions",ascending=False)
```

Out[12]:

	ID	Runtime_minutes	Rating	Votes	Revenue_millions
Year					
2016	111804	31890	1911.7	14431751	11211.65
2015	62407	14541	838.5	14697230	8854.12
2014	50272	11220	670.1	19985162	7997.40
2013	50017	10700	610.0	10000710	7000.70

2013	50617	10562	619.9	19933518	7666.72
ID	Runtime_minutes	Rating	Votes	Revenue_millions	
2012	34104	7623	443.2	18254470	6910.29
Year					
2010	35239	6668	409.6	15166939	5989.65
2011	37143	7220	430.8	15169789	5431.96
2009	29928	5922	355.0	13044813	5292.26
2008	32533	5763	352.8	14326280	5053.22
2007	30559	6446	378.1	12949545	4306.23
2006	25894	5317	313.5	11848758	3624.46

In [13]:

```
# What is the maximum revenue out of (10,20,30,40,50) rows?

df
x = df.iloc[10:51:10]
x1 = x["Revenue_millions"].max()
x1
```

Out[13]:

936.63

In [14]:

```
# How many movies with the genres 'Adventure', 'Action', 'Horror', and
# 'Crime' exist in the IMDB dataset?

df[(df["Genre"].isin(["Adventure", 'Action', "Horror", "Crime"]))].count()
```

Out[14]:

ID 485
Title 485
Genre 485
Director 485
Year 485
Runtime\_minutes 485
Rating 485
Votes 485
Revenue\_millions 436
dtype: int64

# REMOVE AND REPLACE NULL VALUES CLEANED IMDB DATASET

In [15]:

```
df = df.dropna(how='any')
df
```

Out[15]:

	ID	Title	Genre	Director	Year	Runtime_minutes	Rating	Votes	Revenue_millions
0	1	Guardians of the Galaxy	Action	James Gunn	2014	121	8.1	757074	333.13
1	2	Prometheus	Adventure	Ridley Scott	2012	124	7.0	485820	126.46
2	3	Split	Horror	M. Night Shyamalan	2016	117	7.3	157606	138.12
3	4	Sing	Animation	Christophe Lourdelet	2016	108	7.2	60545	270.32
4	5	Suicide Squad	Action	David Ayer	2016	123	6.2	393727	325.02
...	...	...	...	...	...	...	...	...	...
993	994	Resident Evil: Afterlife	Action	Paul W.S. Anderson	2010	97	5.9	140000	60.13

	ID	Title	Genre	Director	Year	Runtime_minutes	Rating	Votes	Revenue_millions
994	995	Project X	Comedy	Nima Nourizadeh	2012	88	6.7	164088	54.72
996	997	Hostel: Part II	Horror	Eli Roth	2007	94	5.5	73152	17.54
997	998	Step Up 2: The Streets	Drama	Jon M. Chu	2008	98	6.2	70699	58.01
999	1000	Nine Lives	Comedy	Barry Sonnenfeld	2016	87	5.3	12435	19.64

872 rows x 9 columns

In [16]:

```
# Create a genre-level report with metrics average rating, the average number
# of votes, and average revenue. What is the average rating of the 'Horror' genre?
# (Round to 2 decimal places)

a = df[(df["Genre"]=="Horror")]
a1 = a["Rating"].mean()
a2 = round(a1,2)
a2
```

Out[16]:

6.1

In [17]:

```
# What is the % revenue of the movie 'Split' in its respective genre and year?

x= df[(df["Title"]=="Split")]
print(x)
x1 = df[(df["Genre"]=="Horror") & (df["Year"]==2016)]
print(x1)
x2 = x1["Revenue_millions"].sum()
print(x2)
perc = x["Revenue_millions"] * 100 / x2
print(perc)
print('\n')
print("% revenue of the movie 'Split' in its respective genre and year :", perc)
```

	ID	Title	Genre	Director	Year	Runtime_minutes	Rating	\
2	3	Split	Horror	M. Night Shyamalan	2016	117	7.3	
	Votes	Revenue_millions						
2	157606	138.12						
	ID	Title	Genre	Director	Year	\		
2	3	Split	Horror	M. Night Shyamalan	2016			
27	28	Dead Awake	Horror	Phillip Guzman	2016			
97	98	The Void	Horror	Jeremy Gillespie	2016			
116	117	The Neon Demon	Horror	Nicolas Winding Refn	2016			
178	179	The Conjuring 2	Horror	James Wan	2016			
193	194	Morgan	Horror	Luke Scott	2016			
258	259	Lights Out	Horror	David F. Sandberg	2016			
461	462	The Boy	Horror	William Brent Bell	2016			
531	532	Friend Request	Horror	Simon Verhoeven	2016			
723	724	Blair Witch	Horror	Adam Wingard	2016			
748	749	Ouija: Origin of Evil	Horror	Mike Flanagan	2016			
776	777	31	Horror	Rob Zombie	2016			
	Runtime_minutes	Rating	Votes	Revenue_millions				
2	117	7.3	157606	138.12				
27	99	4.7	523	0.01				
97	90	5.8	9247	0.15				
116	118	6.2	50359	1.33				
178	134	7.4	137203	102.46				
193	92	5.8	22107	3.91				
258	81	6.4	69823	67.24				
461	97	6.0	51235	35.79				
531	92	5.4	12758	64.03				
723	89	5.1	26088	20.75				
748	99	6.1	30035	34.90				
776	102	5.1	10871	0.78				

469.4699999999999

2 29.42041

Name: Revenue\_millions, dtype: float64

% revenue of the movie 'Split' in its respective genre and year : 2 29.42041

Name: Revenue\_millions, dtype: float64

In [18]:

```
# Add a column 'Votes_norm' in the IMDB dataset using apply() function where Votes_norm is
# [Votes - min(Votes)]*10/[max(votes) - min(votes)]
# The above formula is the normalization formula and converts votes into a scale of 0-10.

# What is the average 'Votes_norm' ? (Round to two decimal places)

df['Votes_norm'] = df['Votes'].apply(lambda x: ((x - df['Votes'].min()) * 10) / (df['Votes'].max() - df['Votes'].min()))
avg_votes = df["Votes_norm"].mean()
round_avg = round(avg_votes, 2)
round_avg
```

Out[18]:

1.06

In [43]:

```
# Add a column 'Total_rating' in the IMDB dataset using apply()
# function where Total_rating is 'Rating' + 'Votes_norm'. What is the highest 'Total_rating' ?

df["Total_rating"] = df["Rating"] + df["Votes_norm"]
highest_Rating = df["Total_rating"].max()
highest_Rating
```

Out[43]:

19.0

In [57]:

```
# How many directors have created movies in the highest number of genres?

# Step 1: Group the dataframe by 'director' and 'genre' columns and count the number of genres each director has created movies in
director_genre_counts = df.groupby(['Director', 'Genre']).size().reset_index(name='count')

# Step 2: Find the maximum count of genres among the directors
max_genre_count = director_genre_counts['count'].max()

# Step 3: Filter the grouped dataframe to select only the directors who have created movies in the highest number of genres

directors_with_max_genres = director_genre_counts[director_genre_counts['count'] == max_genre_count]['Director']

num_directors_highest_genres = directors_with_max_genres.nunique()

num_directors_highest_genres
```

Out[57]:

1

In [ ]:

